

XIAOCHUANG HAN (Han)

xiaochuang.han@gmail.com ♦ 404-988-6956 ♦ <https://xhan77.github.io/>

EDUCATION

University of Washington

September 2021 - Present

Ph.D. in Computer Science and Engineering

Advisor: Yulia Tsvetkov

Carnegie Mellon University

August 2019 - August 2021

M.S. in Language Technologies

Advisor: Yulia Tsvetkov

Georgia Institute of Technology

August 2015 - May 2019

B.S. in Computer Science, Minor in Mathematics

Advisor: Jacob Eisenstein

PUBLICATIONS

Influence Tuning: Demoting Spurious Correlations via Instance Attribution and Instance-Driven Updates

Xiaochuang Han and Yulia Tsvetkov.

Findings of EMNLP 2021

Fortifying Toxic Speech Detectors Against Veiled Toxicity

Xiaochuang Han and Yulia Tsvetkov.

EMNLP 2020

Explaining Black Box Predictions and Unveiling Data Artifacts through Influence Functions

Xiaochuang Han, Byron C. Wallace, and Yulia Tsvetkov.

ACL 2020

Unsupervised Domain Adaptation of Contextualized Embeddings for Sequence Labeling

Xiaochuang Han and Jacob Eisenstein.

EMNLP 2019

No Permanent Friends or Enemies: Tracking Dynamic Relationships between Nations from News

Xiaochuang Han, Eunsol Choi, and Chenhao Tan.

NAACL 2019

Mind Your POV: Convergence of Articles and Editors Towards Wikipedia's Neutrality Norm

Umashanthi Pavalanathan, **Xiaochuang Han**, and Jacob Eisenstein.

CSCW 2018

Interactional Stancetaking in Online Forums

Scott Kiesling, Umashanthi Pavalanathan, Jim Fitzpatrick, **Xiaochuang Han**, and Jacob Eisenstein.

Computational Linguistics, September 2018

Predicting the Suitability of Service Animals Using Instrumented Dog Toys

Ceara Byrne, Jay Zuerndorfer, Larry Freil, **Xiaochuang Han**, Andrew Sirolly, Scott Cilliland, Thad

Starner, and Melody Jackson.
IMWUT, December 2017

RESEARCH EXPERIENCE

CMU LTI / UW CSE TsvetShop

August 2019 - Present

Graduate Research Assistant, with Prof. Yulia Tsvetkov

- Demoting spurious correlations in models by instance attribution and instance-driven updates.
- Fortify toxic language classifiers against veiled toxicity using interpretable ML methods.
- Explore the interpretability of NLP models through the lens of training examples.
- Build data and models to forecast scientific mispractice in medical papers by their preregistrations.

Georgia Tech Computational Linguistics Lab

August 2017 - May 2019

Undergraduate Research Assistant, with Prof. Jacob Eisenstein

- Improved unsupervised domain adaptation of contextualized embeddings for sequence labeling.
- Explored variational methods for geo-entity resolution.
- Analyzed the effect of Wikipedia's neutrality norm.
- Worked on stance classifiers in a quantitative model of stancetaking in online forums.

University of Colorado Boulder NLP and CSS Lab

May 2018 - August 2018

Research Intern, with Prof. Chenhao Tan

- Built an unsupervised model to explore entity-to-entity relations in world news.

TEACHING ASSISTANTSHIPS

CMU 11-711: Algorithms for NLP

Fall 2020

Graduate Teaching Assistant, with Prof. Emma Strubell, Yulia Tsvetkov, and Robert Frederking

- Adapted and redesigned homework assignments, gave a lecture on natural language inference and interpretability in neural NLP, led recitations and hosted office hours.

ACADEMIC SERVICE

Reviewer (*outstanding reviewer)

- ICLR 2022, ACL 2021*, NAACL 2021, EACL 2021*, EMNLP 2020, W-NUT 2020

STUDENT ORGANIZATIONS

Georgia Tech Big Data Club

August 2015 - May 2019

President and Lecturer

- Organized weekly meetings and gave lectures on big data and machine learning tools and algorithms.